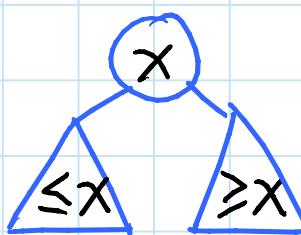
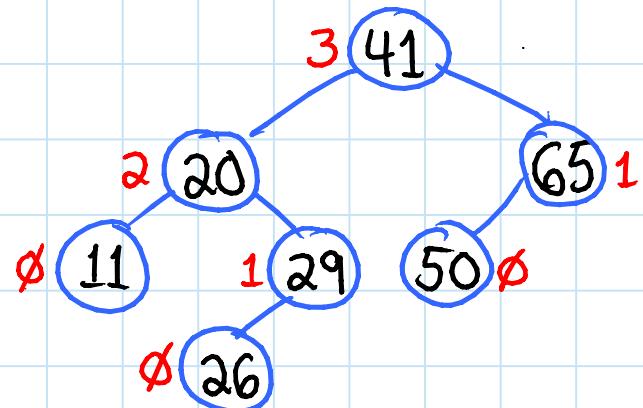


TODAY: Balanced BSTs

- The importance of being balanced
- AVL trees
 - definition & balance
 - rotations
 - insert, delete
- Other balanced trees
- Data structures in general

Recall: Binary Search Trees (BSTs)

- rooted binary tree
- each node has
 - key
 - left pointer
 - right pointer
 - parent pointer
- BST property:

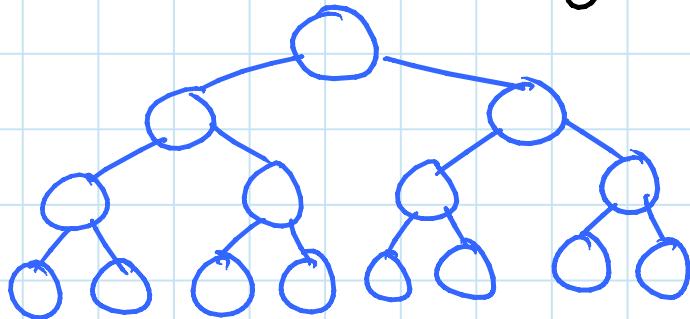


CLRS B.5

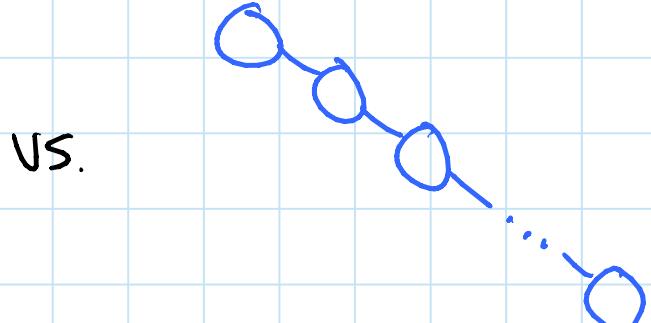
- height of node = length (# edges) of longest downward path to a leaf
- essentially, a more dynamic/flexible version of binary search

The importance of being balanced:

- BSTs support insert, delete, min, max, predecessor, successor, etc. in $O(h)$ time, where $h =$ height of tree
 $(=$ height of root $)$
- h is between $\lg n$ and n :



perfectly balanced



vs.

path

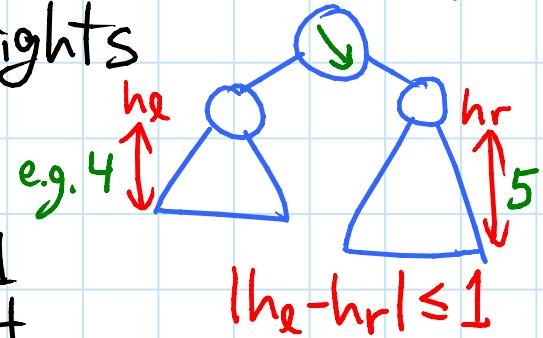
- balanced BST maintains $h = O(\lg n)$
 \Rightarrow all operations run in $O(\lg n)$ time
- today we'll see the first ever such BST data structure, and still one of the simplest: AVL trees

AVL trees: [Adel'son-Vel'skii & Landis 1962]

for every node, require heights of left & right children to differ by at most ± 1

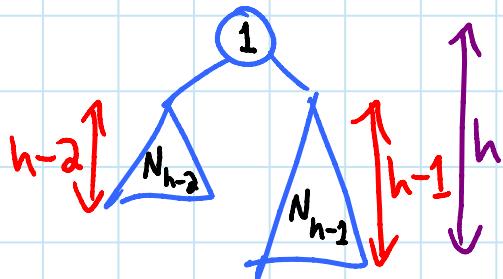
- treat nil tree as height -1
- each node stores its height

(DATA STRUCTURE AUGMENTATION) (like subtree size)
 (Alternatively, can just store difference in heights)



Balance: worst when every node differs by 1

$$\begin{aligned}
 &\text{let } N_h = (\min.) \# \text{ nodes in height-}h \text{ AVL tree} \\
 \Rightarrow N_h &= N_{h-1} + N_{h-2} + 1 \\
 &> 2N_{h-2} \\
 \Rightarrow N_h &> 2^{h-2} \\
 \Rightarrow h &< 2 \lg N_h
 \end{aligned}$$



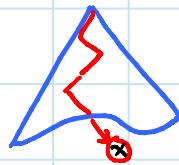
Alternatively: $N_h > F_h$ (n th Fibonacci number)

- in fact $N_h = F_{h+2} - 1$ (simple induction)
- $F_h = \varphi^h / \sqrt{5}$ rounded to nearest integer where $\varphi = \frac{1+\sqrt{5}}{2} \approx 1.618$ (golden ratio)

$$\Rightarrow \max. h \approx \log_{\varphi} n \approx 1.440 \lg n$$

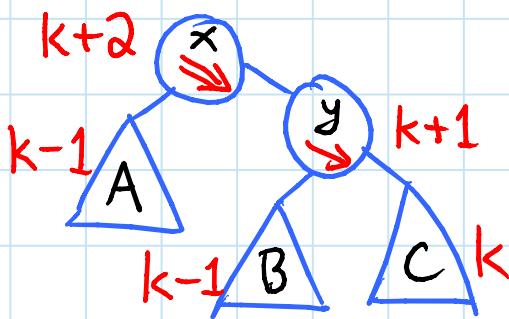
AVL insert:

- ① insert as in simple BST [L5]
- ② walk up tree, restoring AVL property
(and updating heights as you go)

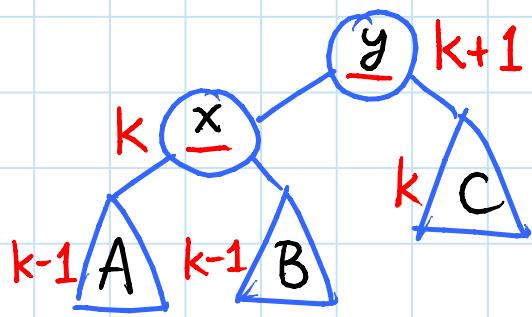


Each step:

- suppose x is lowest node violating AVL
- assume x is right-heavy (left case symmetric)
- if x 's right child is right-heavy or balanced:

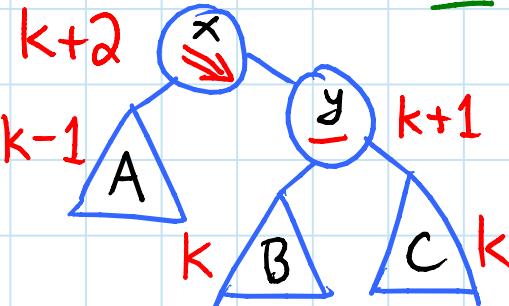


Left-Rotate(x)

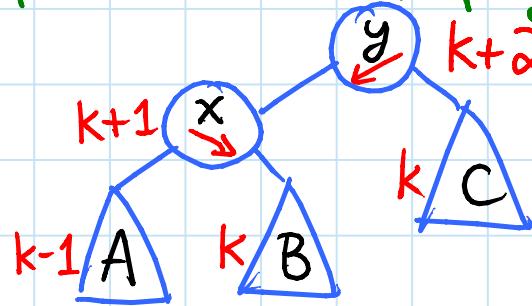


KEY: rotations preserve BST property

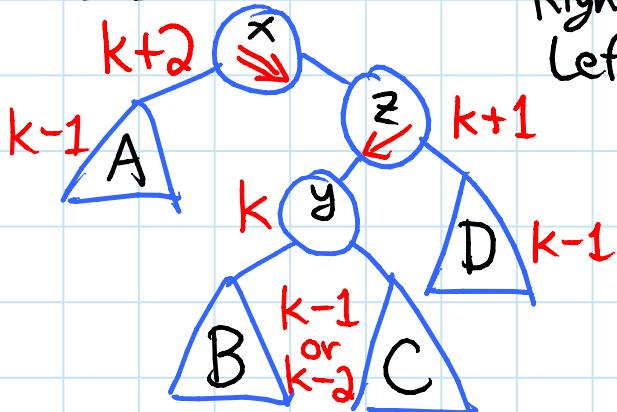
can only happen during Delete



Left-Rotate(x)

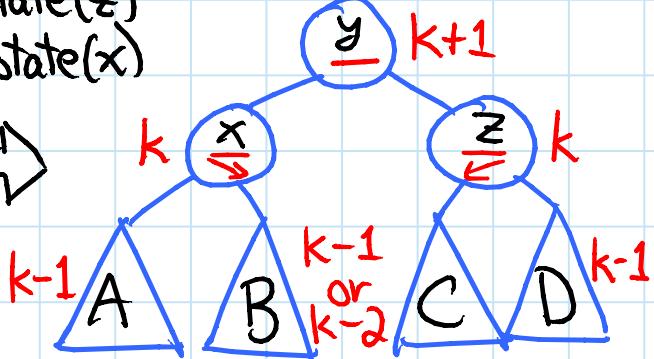


- else:



Right-Rotate(z)
Left-Rotate(x)

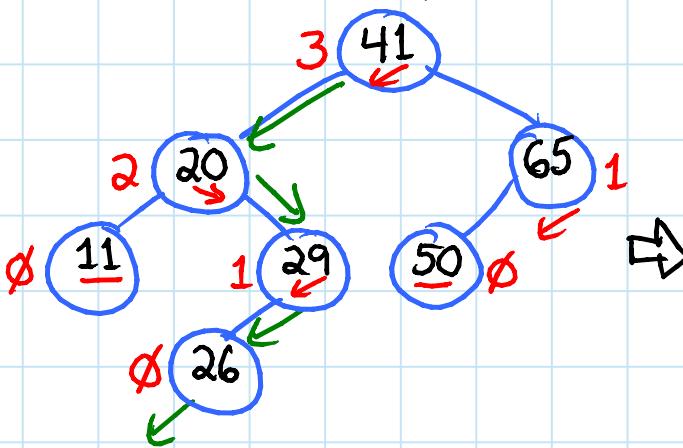
→



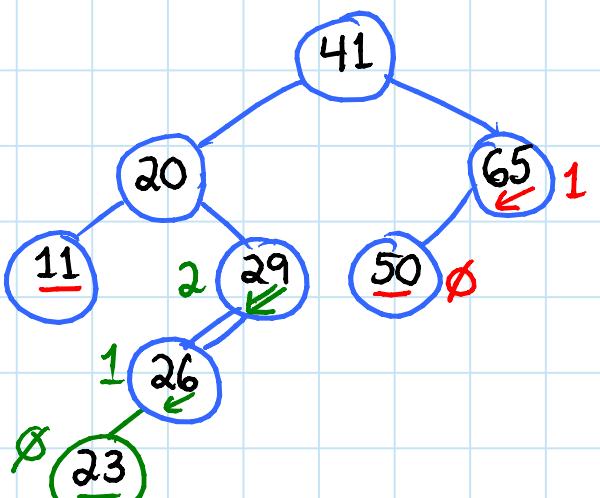
[- then continue up to x 's grandparent, greatgp,...]

Example:

Insert(23)

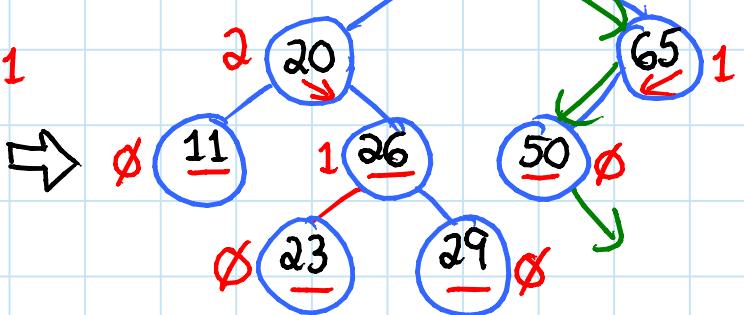
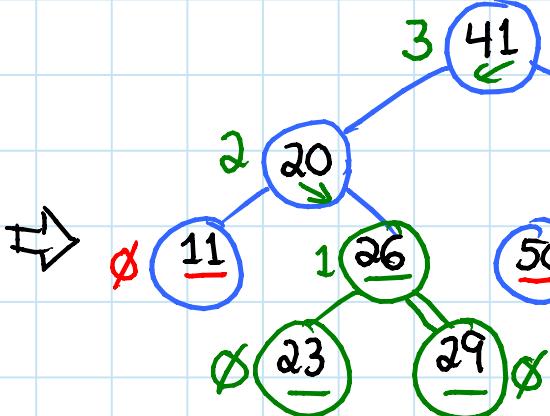


$x=23$: left-left case

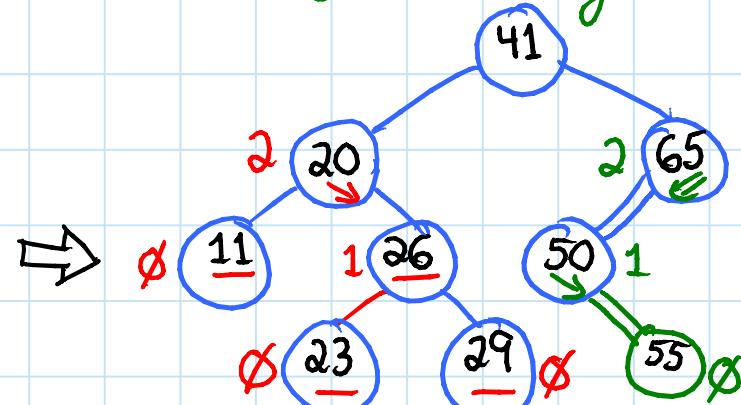


Done.

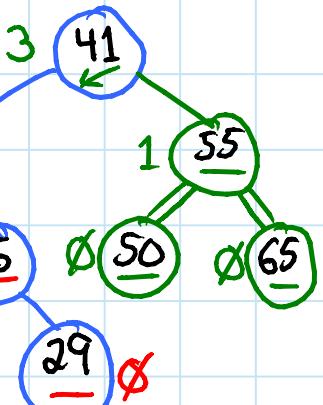
Insert(55)



$x=55$: left-right case



Done.



- in fact never need > 1 rotation pair for an Insert:
height of subtree after rebalancing =
height of subtree before insert

Delete is similar: (but may need $\Theta(\lg n)$ rotations)

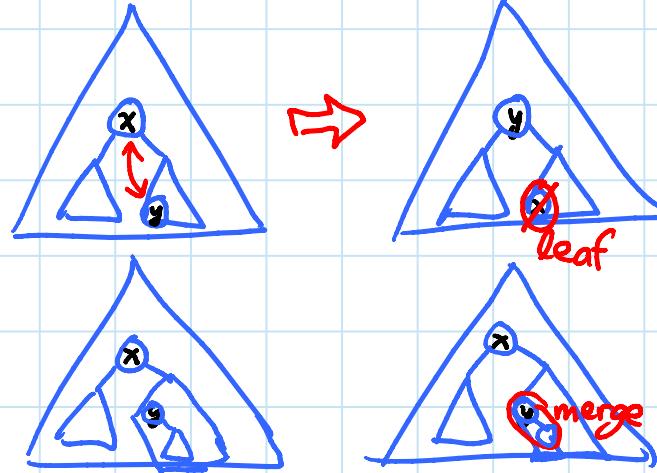
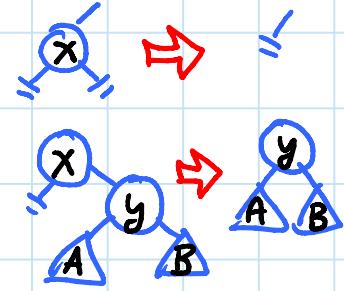
① delete as in simple BST

- if x is a leaf: delete it

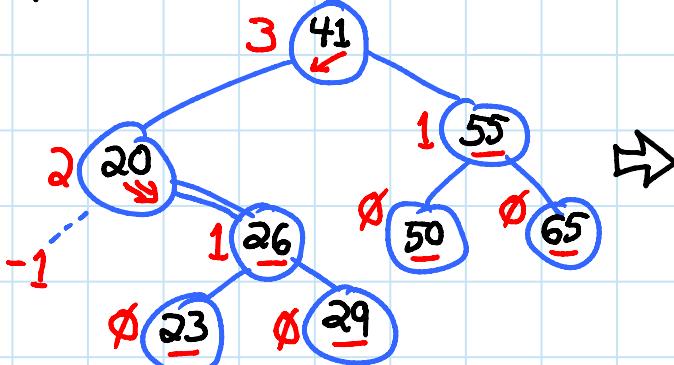
- else if x has exactly 1 child: merge x with its child

- else swap x with its successor, then apply previous case

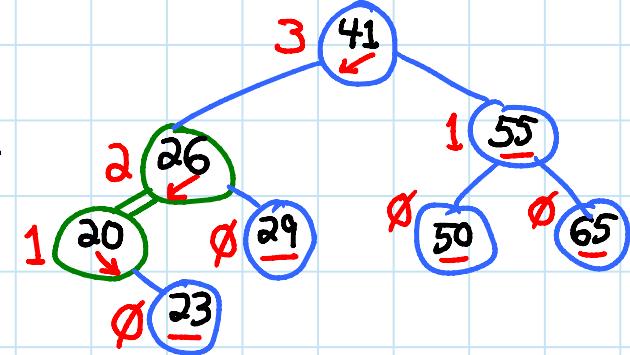
② walk up tree, restoring AVL property (& updating heights)



Example: Delete(11)



$x=20$: right-flat case



AVL Sort:

- insert each item into AVL tree
- in-order traversal

$\Theta(n \lg n)$
 $\Theta(n)$
 $\underline{\Theta(n \lg n)}$

Balanced search trees: there are many!

- AVL trees
- B-trees / 2-3-4 trees
- BB $[\alpha]$ trees
- red-black trees
- (A) - splay trees
- (R) - skip lists
- (A) - scapegoat trees
- (R) - treaps

[Adel'son-Velskii & Landis 1962]
[Bayer & McCreight 1972] [CLRS 18]
[Nievergelt & Reingold 1973]
[CLRS ch. 13]
[Sleator & Tarjan 1985]
[Pugh 1989]
[Galperin & Rivest 1993]
[Seidel & Aragon 1996]

(R) = use random numbers to make decisions
fast with high probability
(A) = "amortized": adding up costs for
several operations \Rightarrow fast on average

e.g. Splay trees are a current research topic
- see 6.854 (Advanced Algorithms)
& 6.851 (Advanced Data Structures)

Big picture:

Abstract Data Type (ADT): interface spec.
 vs. Data Structure (DS): algorithm for each op.

- many possible DSs for one ADT

Priority Queue ADT:

- $Q = \text{new-empty-queue}()$
- $Q.\text{insert}(x)$
- $x = Q.\text{deletemin}()$
- $x = Q.\text{findmin}()$

heap
 $\Theta(1)$
 $\Theta(\lg n)$
 $\Theta(\lg n)$
 $\Theta(1)$

AVL tree
 $\Theta(1)$
 $\Theta(\lg n)$
 $\Theta(\lg n)$
 $\Theta(\lg n)$
 $\hookleftarrow \Theta(1)$

Predecessor/Successor ADT:

- $S = \text{new-empty}()$
- $S.\text{insert}(x)$
- $S.\text{delete}(x)$
- $y = S.\text{predecessor}(x)$
- $y = S.\text{successor}(x)$

heap
 $\Theta(1)$
 $\Theta(\lg n)$
 $\Theta(\lg n)$
 $\Theta(n)$
 $\Theta(n)$

AVL tree
 $\Theta(1)$
 $\Theta(\lg n)$
 $\Theta(\lg n)$
 $\Theta(\lg n)$
 $\Theta(\lg n)$